

Predict__413__Sec55__Homework__1

Singh, Gurjeet

January 28, 2018

Contents

Chapter 2	2
Question 1	2
Ch2.Q1.a)	2
Ch2.Q1.b)	4
Ch2.Q1.c)	4
Question 2	6
Ch2.Q2.a)	6
Ch2.Q2.b)	6
Ch2.Q2.c)	7
Ch2.Q2.d)	8
Question 3	10
Ch2.Q3.a)	10
Ch2.Q3.b&c)	10
Question 4	12
Ch2.Q4.a)	12
Ch2.Q4.b&c)	13
Chapter 4	15
Question 1	15
Ch4.Q1.a)	15
Ch4.Q1.b)	16
Ch4.Q1.c)	18
Ch4.Q1.d)	19
Question 2	19
Ch4.Q2.a)	19
Ch4.Q2.b)	20
Ch4.Q2.c)	21
Ch4.Q2.d)	21
Ch4.Q2.e&f)	23
Question 3	24
Chapter 6	24
Question 1	24
Question 2	24
Ch6.Q2.a)	24
Ch6.Q2.b&c)	25
Ch6.Q2.d)	26
Ch6.Q2.e)	27
Ch6.Q2.f)	29
Ch6.Q2.g)	30
Ch6.Q2.h)	31
Question 3	32
Ch6.Q3.a)	32
Ch6.Q3.b)	32

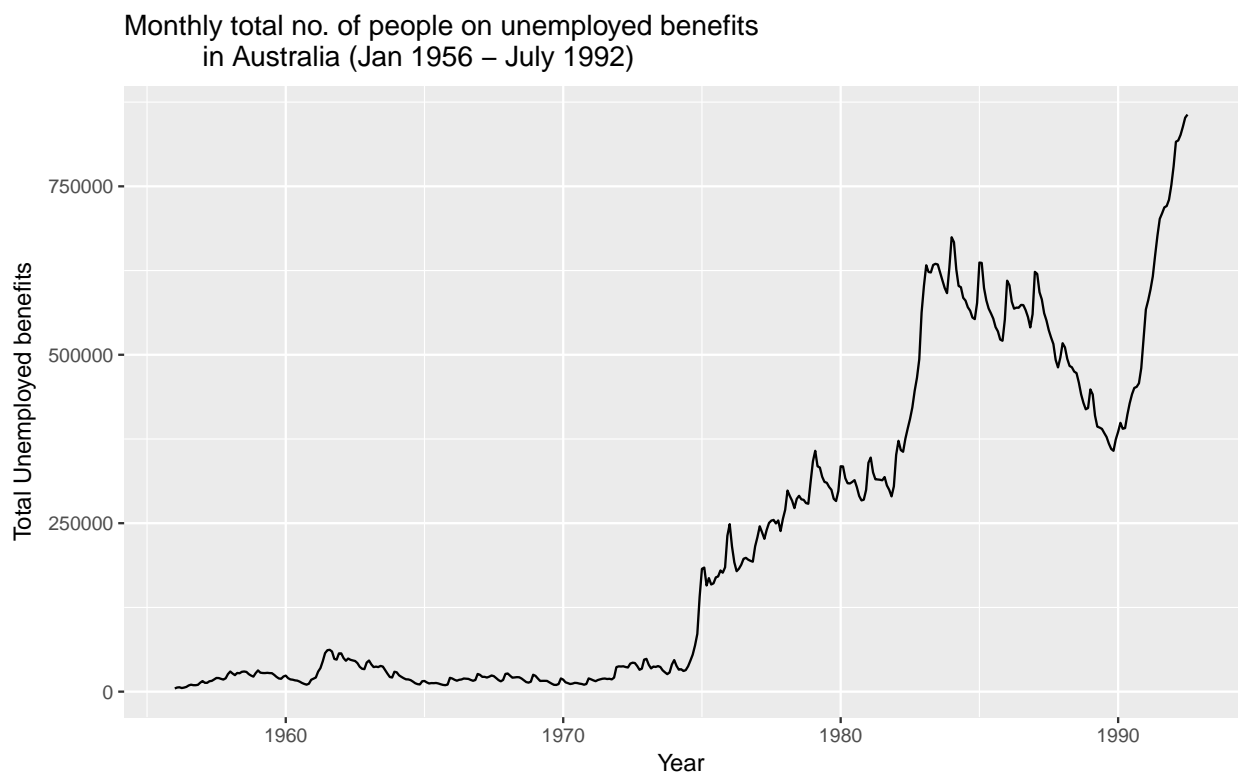
Chapter 2

Question 1

Ch2.Q1.a)


The figure below shows the plot of the monthly total of people on unemployment benefits in Australia from January 1956 to July 1992.

```
data("dole")
autoplot(dole) +
  ggtitle("Monthly total no. of people on unemployed benefits  
in Australia (Jan 1956 - July 1992)") +
  xlab("Year") + ylab("Total Unemployed benefits")
```



The next figure shows the plot with the BoxCox transformation with the 0.5 parameters. The plot seems to smooth a little after the transformation. It is not easy to identify the pattern between the years 1956 and 1975.


```
autoplot(BoxCox(dole,0.5)) +  
  ggtitle("Monthly total no. of people on unemployed benefits  
    in Australia (Jan 1956 - July 1992)") +  
  xlab("Year") + ylab("Transformed total # of unemployed benefits")
```



Ch2.Q1.b)

The figure below shows the plot of the monthly total of accidental deaths in the United States from January 1973 to December 1978.

```
data("usdeaths")
autoplot(usdeaths) +
  ggtitle("Monthly total no. accidental deaths in the
          United States (January 1973–December 1978)") +
  xlab("Year") + ylab("US Deaths")
```




```
# autoplot(BoxCox(usdeaths,0.5)) +
#   ggtitle("Monthly total no. accidental deaths in the
#           United States (January 1973–December 1978)") +
#   xlab("Year") + ylab("US Deaths")
```

Ch2.Q1.c)

The figure below shows the plot of the quarterly production of bricks (in millions of units) at Portland, Australia from March 1956 to September 1994.

```
data("bricksq")
# head(bricksq)
autoplot(bricksq) +
  ggtitle("Quarterly production of bricks (in millions of units)
          at Portland, Australia (March 1956–September 1994)") +
  xlab("Year") + ylab("Bricks (in millions of units)")
```

Quarterly production of bricks (in millions of units)
at Portland, Australia (March 1956–September 1994)




```
# autoplot(BoxCox(bricksq, 0.5)) +  
# ggtitle("Quarterly production of bricks (in millions of units)  
#         at Portland, Australia (March 1956–September 1994)") +  
# xlab("Year") + ylab("Bricks (in millions of units)")
```

Question 2

Ch2.Q2.a)

The figure below shows the time series plot of Dow Jones index from August 28, 1972 to December 18, 1972.


```
data("dowjones")
autoplot(dowjones) + xlab("Day") + ylab("Dow Jones Index") +
  ggtitle("Dow-Jones index, 28 Aug - 18 Dec 1972")
```



Ch2.Q2.b)

The figure below shows the forecast of the Dow Jones Index data using the drift method.

```
autoplot(rwf(dowjones, drift=TRUE, h=20)) +
  forecast::autolayer(rwf(dowjones, drift=TRUE, h=20), PI=FALSE, series="Drift") +
  ggtitle("Dow-Jones index, 28 Aug - 18 Dec 1972") +
  xlab("Day") + ylab("Dow Jones Index") +
  guides(colour=guide_legend(title="Forecast"))
```




Ch2.Q2.c)

The figure below shows that the forecasts are identical to extending the line drawn between the first and last observations.

```
Xk <- length(dowjones)
X1 <- 1
Yk <- dowjones[Xk]
Y1 <- dowjones[1]

slope <- (Yk - Y1)/(Xk - X1)
intercept <- Y1 - slope


autoplot(rwf(dowjones, drift=TRUE, h=20)) +
  forecast::autolayer(rwf(dowjones, drift=TRUE, h=20), PI=FALSE, series="Drift") +
  ggtitle("Dow-Jones index, 28 Aug - 18 Dec 1972") +
  xlab("Day") + ylab("Dow Jones Index") +
  guides(colour=guide_legend(title="Forecast")) +
  geom_abline(slope = slope, intercept = intercept, col = "khaki4")
```



Ch2.Q2.d)

The figure below shows some of the benchmark functions such as Mean, Drift, and Naive methods to forecast the Dow Jones Index data. It appears that drift method is the clear winner here.

```
autoplot(dowjones) +
  forecast::autolayer(meanf(dowjones, h=40), PI=FALSE, series="Mean") +
  forecast::autolayer(rwf(dowjones, h=40), PI=FALSE, series="Na?ve") +
  forecast::autolayer(rwf(dowjones, drift=TRUE, h=40), PI=FALSE, series="Drift") +
  ggtitle("Dow-Jones index, 28 Aug - 18 Dec 1972") +
  xlab("Day") + ylab("Dow Jones Index") +
  guides(colour=guide_legend(title="Forecast"))
```

The table below shows the metrics from each benchmark functions used to forecast the data. Looking at the results, it is quite evident that drift method performed really well and is the best method to select. However, the Naive method did not perform so poorly either. It appears that Naive method prediction is within 95% prediction intervals. Therefore, the metric results are very close to what drift method.

```
meanfcst <- accuracy(meanf(dowjones, h=40))
row.names(meanfcst) <- "Mean"
naivefcst <- accuracy(rwf(dowjones, h=40))
row.names(naivefcst) <- "Na?ve"
driftfcst <- accuracy(rwf(dowjones, drift=TRUE, h=40))
row.names(driftfcst) <- "Drift"

knitr::kable(rbind(meanfcst, naivefcst, driftfcst), caption = "Benchmark Results")
```

Table 1: Benchmark Results

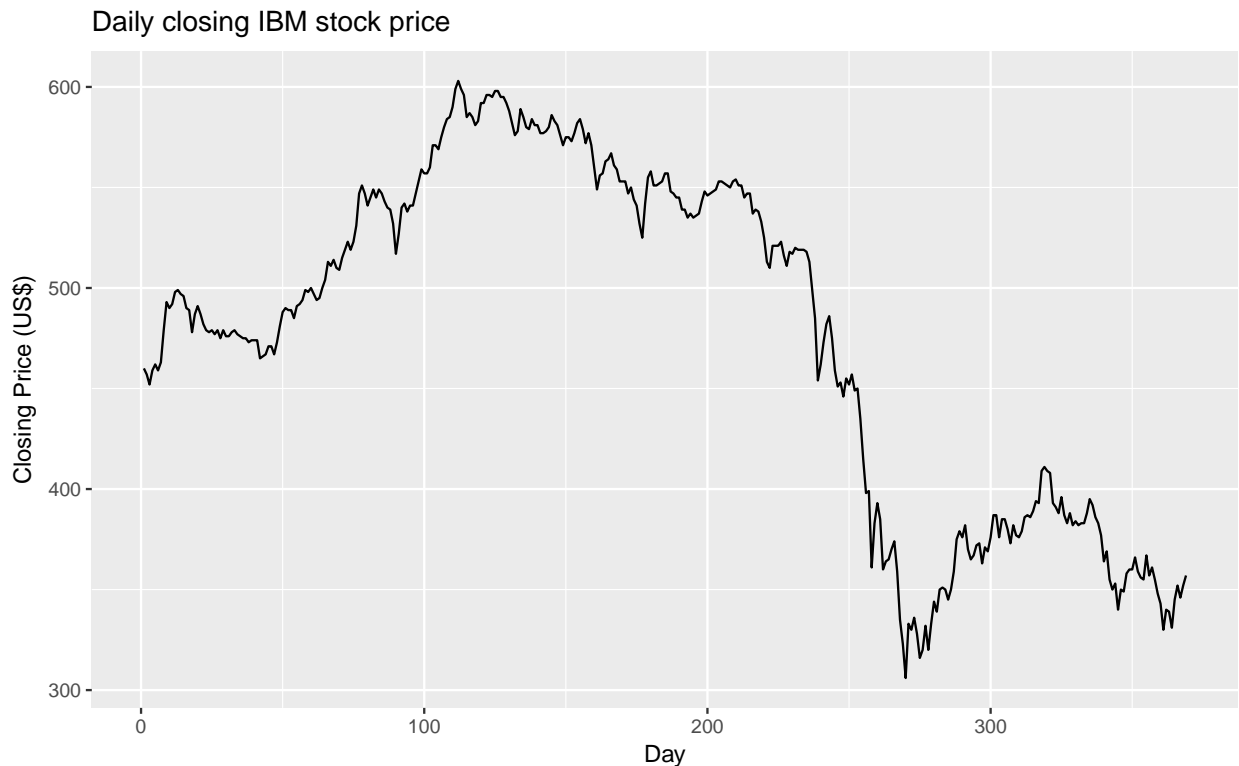
	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Mean	0.0000000	5.4706531	5.1041026	-0.2216950	4.3998889	14.937890	0.9853294
Na?ve	0.1336364	0.4447223	0.3416883	0.1144757	0.2936792	1.000000	0.4218786
Drift	0.0000000	0.4241689	0.3253365	-0.0012387	0.2796548	0.952144	0.4218786

Question 3

Ch2.Q3.a)

The figure below shows the time series plot of the daily closing IBM stock prices.

```
data("ibmclose")
autoplot(ibmclose) + xlab("Day") + ylab("Closing Price (US$)") +
  ggtitle("Daily closing IBM stock price")
```




Ch2.Q3.b&c)

The figure below shows various benchmark functions such as Mean, Drift, and Naive methods used to forecast the daily closing IBM stock prices. It appears that naive and drift methods performed well here. They are well within the 80 and 95% prediction interval.

```
#str(ibmclose)
train.IBM.ts <- window(ibmclose, end = 300)
test.IBM.ts <- window(ibmclose, start = 301, end = 369)
# str(train.IBM.ts)
# str(test.IBM.ts)
Fcast.IBM.mean <- meanf(train.IBM.ts, h=40)
Fcast.IBM.Naive <- rwf(train.IBM.ts, h=40)
Fcast.IBM.Drift <- rwf(train.IBM.ts, drift=TRUE, h=40)

autoplot(train.IBM.ts) +
  forecast::autolayer(Fcast.IBM.mean, PI=FALSE, series="Mean") +
```

```
forecast::autolayer(Fcast.IBM.Naive, PI=FALSE, series="Na?ve") +
forecast::autolayer(Fcast.IBM.Drift, PI=FALSE, series="Drift") +
ggtitle("Daily closing IBM stock price") +
xlab("Day") + ylab("Closing Price (US$)") +
guides(colour=guide_legend(title="Forecast"))
```



The tables below list all the metrics from each benchmark functions used to forecast the data. Looking at the results, it appears that naive method is the best option here. MAPE and RMSE for ‘Test Set’ for naive method beats the other two.

```
acc.Mean <- accuracy(Fcast.IBM.mean, test.IBM.ts)
acc.Naive <- accuracy(Fcast.IBM.Naive, test.IBM.ts)
acc.Drift <-accuracy(Fcast.IBM.Drift, test.IBM.ts)

knitr::kable(acc.Mean[,1:6], caption = "Metrics for Mean Method")
```

Table 2: Metrics for Mean Method

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	0.0000	73.61532	58.72231	-2.642058	13.03019	11.52098
Test set	-115.7933	116.19611	115.79333	-29.992869	29.99287	22.71798

```
knitr::kable(acc.Naive[,1:6], caption = "Metrics for Naive Method")
```

Table 3: Metrics for Naive Method

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	-0.2809365	7.302815	5.09699	-0.0826287	1.115844	1.000000
Test set	11.1000000	14.719035	11.85000	2.8077817	3.012831	2.324902

```
knitr::kable(acc.Drift[,1:6], caption = "Metrics for Drift Method")
```

Table 4: Metrics for Drift Method

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	0.0000	7.297409	5.127996	-0.0253012	1.121650	1.006083
Test set	16.8592	19.733866	16.948997	4.2961173	4.320445	3.325295


Question 4

Ch2.Q4.a)

The figure below shows the time series plot of the sales of new one-family houses in the USA from January 1973 to November 1995.

```
data("hsales")
autoplot(hsales) + xlab("Year") + ylab("Monthly Sales (US$)") +
  ggtitle("Monthly sales of new one-family houses sold in the USA since 1973")
```

Monthly sales of new one-family houses sold in the USA since 1973




Ch2.Q4.b&c)

The figure below shows various benchmark functions such as Mean, Drift, and Naive methods used to forecast the sales of new one-family houses in the USA from January 1973 to November 1995. It appears that all three methods forecasted really close. It would be ideal to see the metrics result of each method to select the best method.

```
#str(hsales)
train.HS.ts <- window(hsales, Start = 1973, end = c(1993,11))
test.HS.ts <- window(hsales, start = c(1993,12))
# str(train.HS.ts)
# str(test.HS.ts)
Fcast.HS.mean <- meanf(train.HS.ts, h=40)
Fcast.HS.Naive <- rwf(train.HS.ts, h=40)
Fcast.HS.Drift <- rwf(train.HS.ts, drift=TRUE, h=40)

autoplot(train.HS.ts) +
  forecast::autolayer(Fcast.HS.mean, PI=FALSE, series="Mean") +
  forecast::autolayer(Fcast.HS.Naive, PI=FALSE, series="Na?ve") +
  forecast::autolayer(Fcast.HS.Drift, PI=FALSE, series="Drift") +
  ggtitle("Monthly sales of new one-family houses sold in the USA since 1973") +
  xlab("Year") + ylab("Monthly Sales (US$)") +
  guides(colour=guide_legend(title="Forecast"))
```

Monthly sales of new one-family houses sold in the USA since 1973



The tables below list all the metrics from each benchmark functions used to forecast the data. Looking at the results, it appears that naive method is the best option here. MAPE and RMSE for ‘Test set’ for naive method beat the other two.

```
acc.Mean <- accuracy(Fcast.HS.mean, test.HS.ts)
acc.Naive <- accuracy(Fcast.HS.Naive, test.HS.ts)
acc.Drift <-accuracy(Fcast.HS.Drift, test.HS.ts)

knitr::kable(acc.Mean[,1:6], caption = "Metrics for Mean Method")
```

Table 5: Metrics for Mean Method

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	0.000000	12.162811	9.532738	-6.144876	20.38306	1.1234341
Test set	3.839475	9.022555	7.561587	4.779121	13.26183	0.8911338

```
knitr::kable(acc.Naive[,1:6], caption = "Metrics for Naive Method")
```

Table 6: Metrics for Naive Method

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	-0.008000	6.301111	5.000000	-0.767457	9.903992	0.5892505
Test set	2.791667	8.628924	7.208333	2.858639	12.849194	0.8495028

```
knitr::kable(acc.Drift[,1:6], caption = "Metrics for Drift Method")
```

Table 7: Metrics for Drift Method

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	0.000000	6.301106	4.999872	-0.7511048	9.903063	0.5892354
Test set	2.891667	8.658795	7.249000	3.0426108	12.901696	0.8542954

Chapter 4


Question 1

Ch4.Q1.a)

There appears to be a negative relationship between temp and Mwh. However, the relationship is not perfect. As the temperature increases, the electricity consumption decreases. The negative relationship exists because of the less use of heater/heat due to the increase in the temperature. Therefore, the electricity consumption decreases.

```
data("econsumption")
ggplot(data = econsumption, mapping = aes(x = temp, y = Mwh)) +
  geom_point() +
  geom_smooth(method = 'lm', se = F) +
```


```
theme_classic()
```



Ch4.Q1.b)

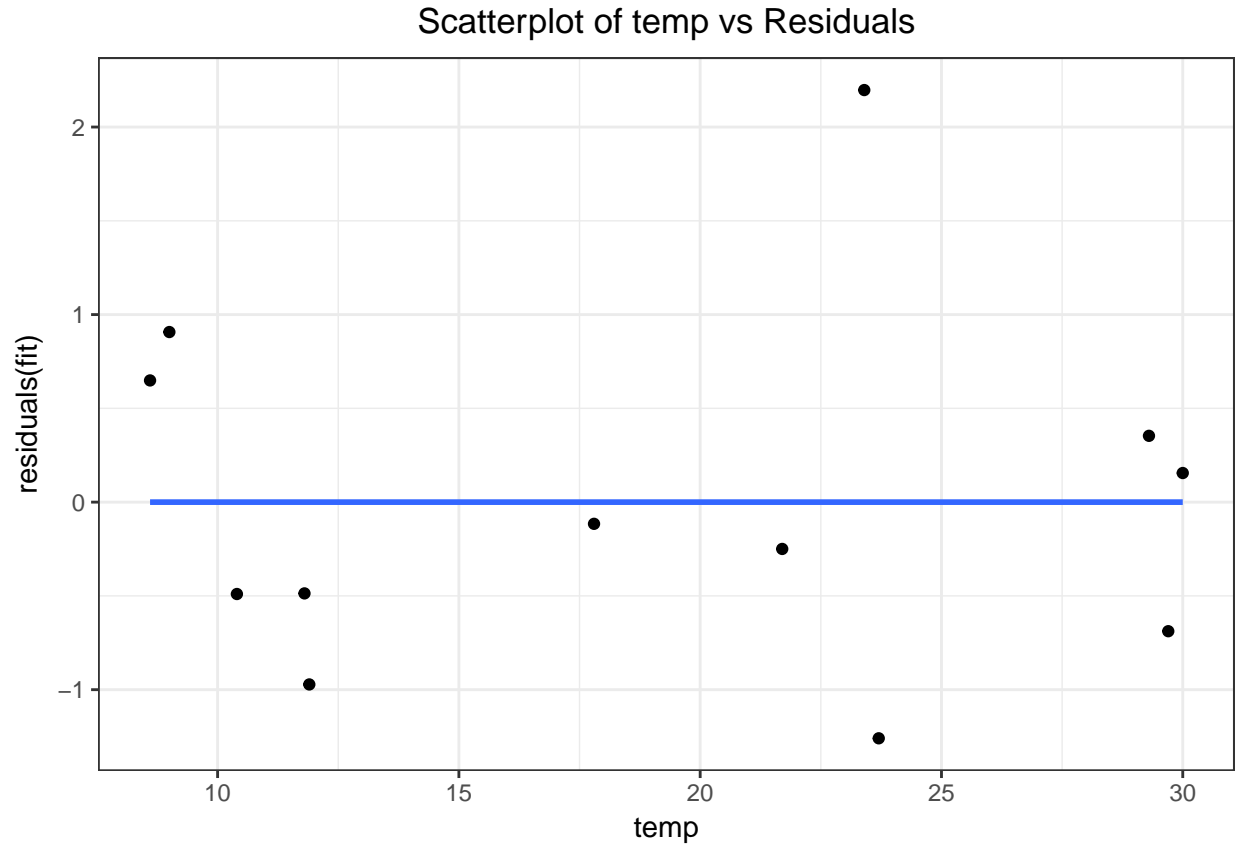
The figure below shows the normal QQ plot. It appears that most of the points fall on the line or close to it. It also shows that there is one outlier present (top right) in the data.

```
fit <- with(econsumption, lm(Mwh ~ temp))
#summary(fit)
#plot(with(econsumption, residuals(fit) ~ temp))
qqnorm(fit$residuals)
qqline(fit$residuals)
```

The figure below shows the residual plots confirming that the errors are random and have constant variance.

```
ggplot(data = econsumption, mapping = aes(x = temp, y = residuals(fit))) +  
  geom_point() +  
  geom_smooth(method = 'lm', se = F) +  
  ggtitle("Scatterplot of temp vs Residuals") +  
  theme_bw() +  
  theme(plot.title = element_text(hjust=0.5))
```



Ch4.Q1.c)

The table below shows the forecasted values. The prediction seems to be accurate. For instance, in the source data, when the temperature was 10.4 degrees, the electricity consumption was 18.2 Mwh. For 10 degrees, our model has predicted 18.75 Mwh which to me seems quite close.

```
tempLab <- data.frame(temp = c(10,35))
fcast <- forecast(fit, newdata = tempLab)
attCol <- cbind.data.frame(tempLab, data.frame(fcast))
knitr::kable(attCol, caption = "Electricity Consumption Prediction
                                when temp = 10 and temp = 35")
```

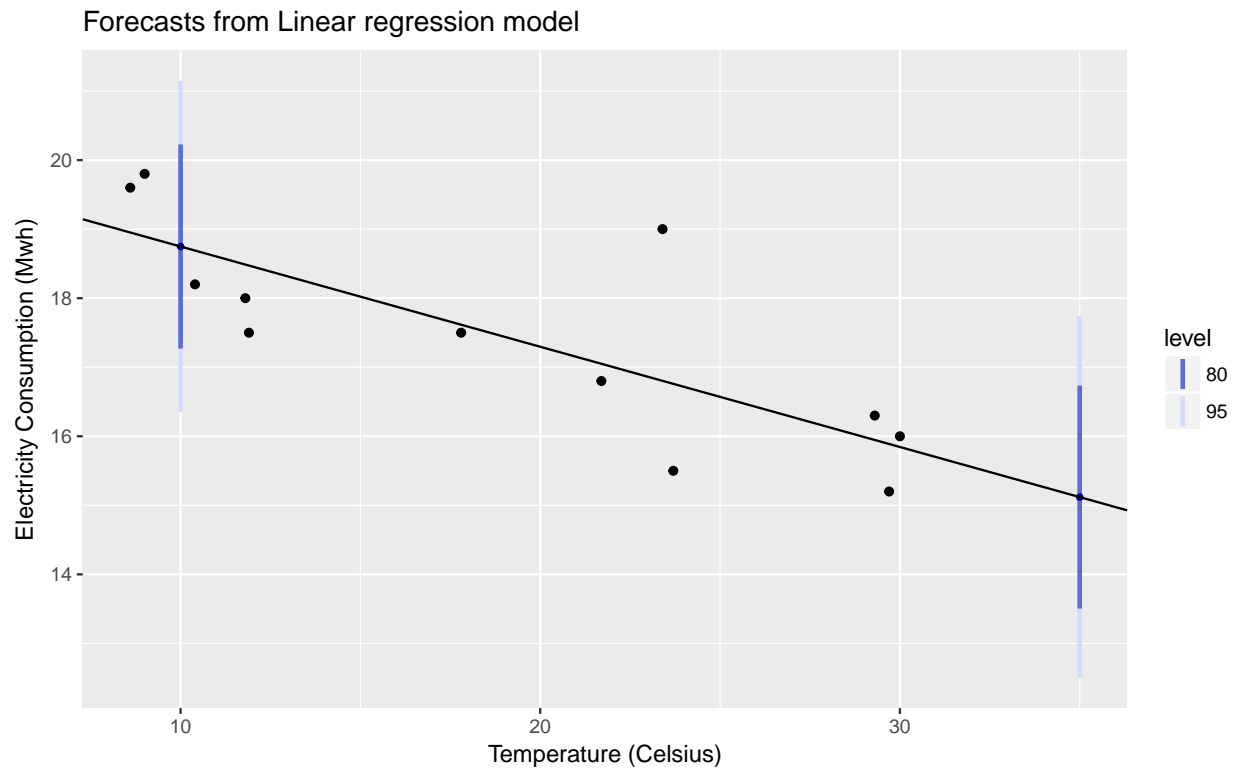
Table 8: Electricity Consumption Prediction when temp = 10 and temp = 35

temp	Point.Forecast	Lo.80	Hi.80	Lo.95	Hi.95
10	18.74795	17.27010	20.22579	16.34824	21.14766
35	15.11902	13.50469	16.73335	12.49768	17.74035

Ch4.Q1.d)

The figure below shows the forecast with 80% and 95% prediction intervals for electricity consumption with temperature = 10 degrees and temperature = 35 degrees.

```
autoplot(fcast) +  
  xlab("Temperature (Celsius)") +  
  ylab("Electricity Consumption (Mwh)")
```



Question 2

Ch4.Q2.a)

The Winning times (in seconds) for men's 400 meters final for each of the last few Olympic games (2000 - 2012) has been included in the original dataset. The table below shows the summary statistic of the new dataset.

```
data("olympic")  
lastfewoly <- data.frame(Year = c(2000, 2004, 2008, 2012),  
  time = c(43.84, 44.00, 43.75, 43.94))  
olympic <- rbind.data.frame(olympic, lastfewoly)  
summ <- cbind(summary(olympic$Year),  
  summary(olympic$time))  
colnames(summ) <- c("Year", "Time")  
knitr::kable(summ,  
  caption = "Summary Statistic")
```

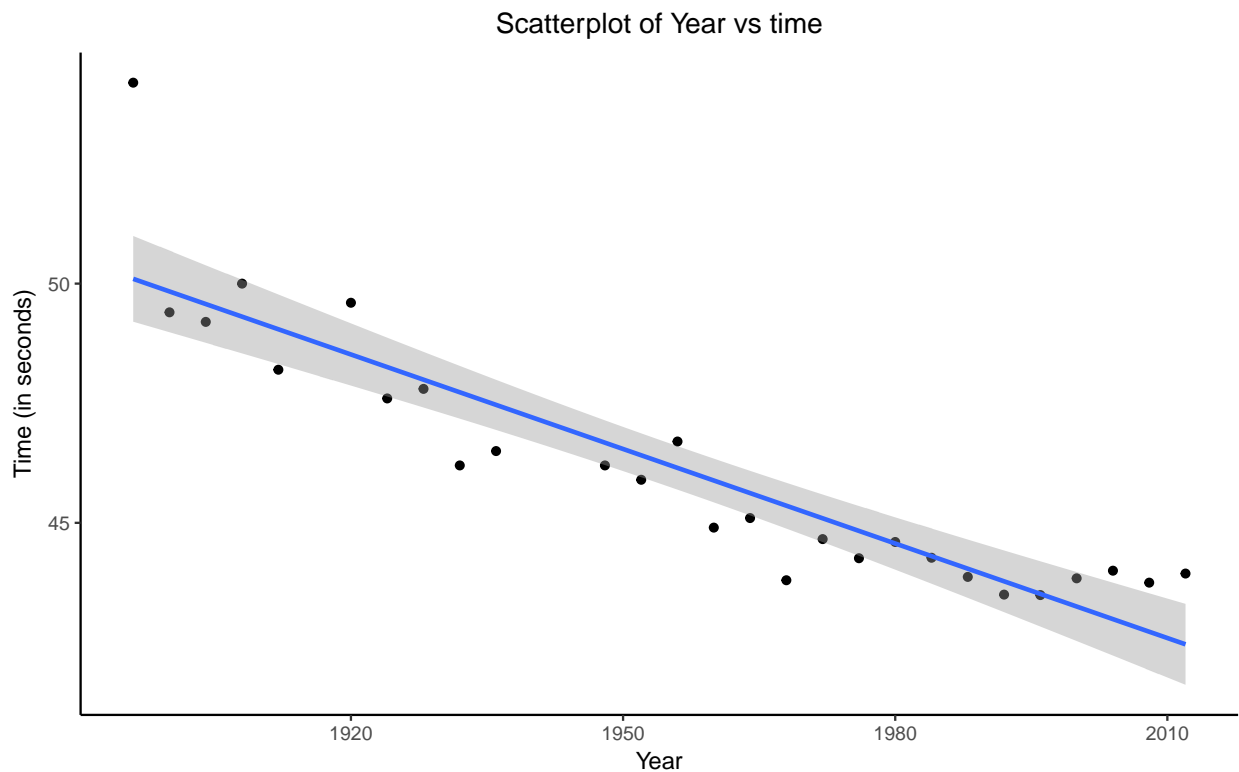
Table 9: Summary Statistic

	Year	Time
Min.	1896	43.49
1st Qu.	1926	43.97
Median	1960	45.10
Mean	1956	46.13
3rd Qu.	1986	47.70
Max.	2012	54.20

Ch4.Q2.b)

The figure below shows the scatterplot of Year vs time with the regression line. The plot shows that the winning times have significantly reduced in the recent years as compared to prior to 1950. There is also an outlier present (top left) in the data.

```
ggplot(data = olympic, mapping = aes(x = Year, y = time)) +
  geom_point() +
  geom_smooth(method = 'lm', se = T) +
  ggtitle("Scatterplot of Year vs time") +
  ylab("Time (in seconds)") +
  theme_classic() +
  theme(plot.title = element_text(hjust=0.5))
```



Ch4.Q2.c)

The result below shows the summary statistic of our linear regression model. The dataset was split into training and test set. We use the 70/30 split approach. We then used the training set to create our regression model. We see that for every one unit increase in the year, the winning times will decrease on average by 0.088.

```
train.df <- olympic[1:18,]
test.df <- olympic[19:27,]
fitReg <- lm(time ~ Year, data = train.df)
summary(fitReg)


##
## Call:
## lm(formula = time ~ Year, data = train.df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.4266 -0.7634 -0.0972  0.4614  3.3965
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 218.11738   20.65829  10.558 1.28e-08 ***
## Year        -0.08825    0.01067  -8.273 3.58e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.161 on 16 degrees of freedom
## Multiple R-squared:  0.8105, Adjusted R-squared:  0.7987
## F-statistic: 68.44 on 1 and 16 DF,  p-value: 3.581e-07
```

Ch4.Q2.d)

The Normal QQ plot below shows that most of the points fall on the line or close to it. It also shows that there are two outliers (one in the top right and one in the bottom left) in the data.

```
qqnorm(fitReg$residuals)
qqline(fitReg$residuals)
```


Normal Q-Q Plot



The residual plot shows that the errors are random. It does seem to me that there is any pattern to it. Hence, we conclude that residuals show constant variance.

```
ggplot(data = train.df, mapping = aes(x = Year, y = residuals(fitReg))) +  
  geom_point() +  
  geom_smooth(method = 'lm', se = F) +  
  ggtitle("Scatterplot of Year vs Residuals") +  
  theme_bw() +  
  theme(plot.title = element_text(hjust=0.5))
```

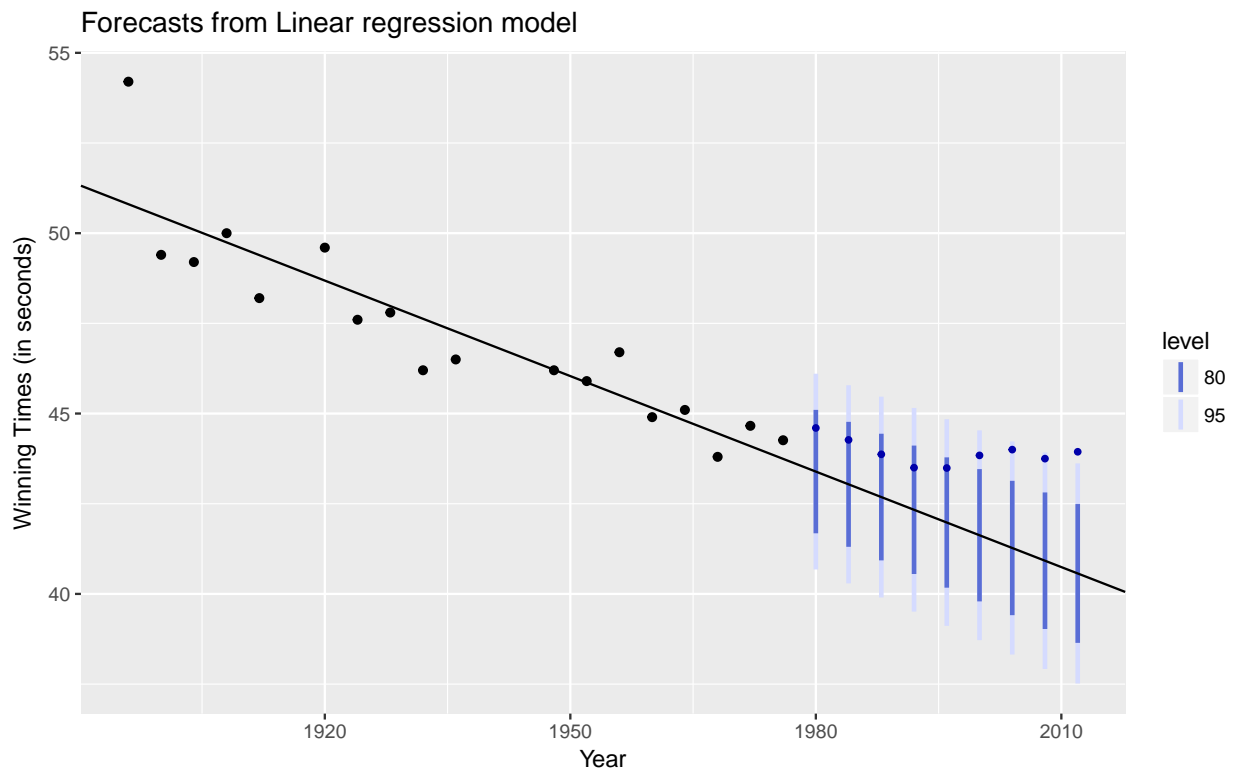
Scatterplot of Year vs Residuals



Ch4.Q2.e&f)

The figure below shows the forecast with 80% and 95% prediction intervals for winning times with year greater than 1976.

```
fcastReg <- forecast(fitReg, newdata = test.df)
YearLab <- data.frame(Year = test.df$Year)
attCol <- cbind.data.frame(YearLab, Obs_Time = test.df$time, data.frame(fcastReg))
autoplot(fcastReg) +
  xlab("Year") +
  ylab("Winning Times (in seconds)")
```



The table below shows the observed winning time and predicted winning time for four Olympics games held in years 2000, 2004, 2008, and 2012. Looks like our model did not project very well as compared to the reality. However, the prediction is within the 95% interval.

```
knitr::kable(attCol[attCol$Year>1996, c('Year','Obs_Time', "Point.Forecast")],
  caption = "Olympic Games Winning
    Times Prediction", row.names = F)
```

Table 10: Olympic Games Winning Times Prediction

Year	Obs_Time	Point.Forecast
2000	43.84	41.62594
2004	44.00	41.27296
2008	43.75	40.91997
2012	43.94	40.56699

Question 3

The steps below show the derivation of the log-log model to prove that coefficient B1 is the elasticity coefficient.

Step 1: $\log y = B_0 + B_1 \log x + e$ (log-log model)

Step 2: $ddy[\log y] = ddy[B_0] + ddy[B_1 * \log x] + ddy[e]$ (taking derivatives on each side)

Step 3: $1/y = 0 + B_1 ddy[\log x] + 0$ (ddy values $ddy[\log y] = 1/y$; $ddy[\log x] = 1/x$; $ddy[B_0] = 0 = ddy[e]$)

Step 4: $1/y = B_1 * 1/x$

Step 5: $x = B_1 y$

Step 6: $x/y = B_1$

Chapter 6

Question 1

The following transformation shows that 3 X 5 MA is equivalent to a 7-term weighted moving average with weights of 0.067, 0.133, 0.200, 0.200, 0.200, 0.133, and 0.067

$$\begin{aligned} &= 1/3[1/5(Y_{t-3} + Y_{t-2} + Y_{t-1} + Y_t + Y_{t+1}) + 1/5(Y_{t-2} + Y_{t-1} + Y_t + Y_{t+1} + Y_{t+2}) + 1/5(Y_{t-1} + Y_t \\ &+ Y_{t+1} + Y_{t+2} + Y_{t+3})] \\ &= 1/3[1/5\{ Y_{t-3} + Y_{t-2} + Y_{t-1} + Y_t + Y_{t+1} + Y_{t-2} + Y_{t-1} + Y_t + Y_{t+1} + Y_{t+2} + Y_{t-1} + Y_t + Y_{t+1} \\ &+ Y_{t+2} + Y_{t+3}\}] \\ &= 1/15[Y_{t-3} + 2(Y_{t-2}) + 3(Y_{t-1}) + 3(Y_t) + 3(Y_{t+1}) + 2(Y_{t+2}) + Y_{t+3}] \\ &= 1/15(Y_{t-3}) + 2/15(Y_{t-2}) + 1/5(Y_{t-1}) + 1/5(Y_t) + 1/5(Y_{t+1}) + 2/15(Y_{t+2}) + 1/15(Y_{t+3}) \\ &= 0.067(Y_{t-3}) + 0.133(Y_{t-2}) + 0.2(Y_{t-1}) + 0.2(Y_t) + 0.2(Y_{t+1}) + 0.133(Y_{t+2}) + 0.067(Y_{t+3}) \end{aligned}$$

Question 2

Ch6.Q2.a)

Yes, we identify the seasonal fluctuations and the trend in the plot below. Every year there's a slight decrease in the sale, in the beginning, sale picks up in the mid of the year, and then starts to decrease again. Also, the sale has been increasing year over year. Hence, there's also an increasing trend in the sale of product A.

```
data("plastics")
autoplot(plastics) +
  ylab("Sale Amount (in thousands)") +
  xlab("Year") +
  ggtitle("Monthly sales of product A
          for a plastics manufacturer")
```





Ch6.Q2.b&c)

The figure below shows trend-cycle and seasonal indices of the classical multiplicative decomposition.

Yes, the results in the figure below support the graphical interpretation of part(a). With this decomposition, we confirm that the seasonal fluctuations and the increasing trend are present in the data.

```
plastics %>%  
  decompose(type="multiplicative") %>%  
  autoplot() + xlab("Year") +  
  ggtitle("Classical multiplicative decomposition of product A")
```

Classical multiplicative decomposition of product A



Ch6.Q2.d)

The figure below shows the plot of the original data and seasonally adjusted data.

```
fit <- stl(plastics, s.window="periodic", robust = T)
autoplot(plastics, series = "Data") +
  forecast::autolayer(seasadj(fit), series="Seasonally Adjusted") +
  xlab("Year") + ylab("Sale Amount (in thousands)") +
  ggtitle("Product A monthly sales with seasonally adjusted") +
```

```
scale_colour_manual(values=c("darkgoldenrod", "blue"),
  breaks=c("Data", "Seasonally Adjusted"))
```

Product A monthly sales with seasonally adjusted



Ch6.Q2.e)

In the original data, we replaced the amount to 2,023 for the month of December and year 3. The tables below show the changes in the seasonally adjusted data. The first table shows the result from the original data set and second table shows the result from the data set with an outlier.

```
#Seasonally adjusted data from original data set
seasFit <- seasadj(fit)
seasFit
```

	Jan	Feb	Mar	Apr	May	Jun	Jul
## 1	995.8392	1014.8402	1013.9941	989.3123	984.0645	949.6119	967.2812
## 2	994.8392	1017.8402	1011.9941	1023.3123	1053.0645	1065.6119	1092.2812
## 3	1149.8392	1110.8402	1122.9941	1146.3123	1158.0645	1168.6119	1105.2812
## 4	1204.8392	1178.8402	1175.9941	1200.3123	1228.0645	1264.6119	1288.2812
## 5	1283.8392	1349.8402	1363.9941	1376.3123	1422.0645	1479.6119	1413.2812
	Aug	Sep	Oct	Nov	Dec		
## 1	949.2518	935.0209	952.1319	995.7861	976.8658		
## 2	1082.2518	1068.0209	1117.1319	1090.7861	1094.8658		
## 3	1169.2518	1200.0209	1274.1319	1194.7861	1216.8658		
## 4	1288.2518	1331.0209	1421.1319	1427.7861	1402.8658		
## 5	1341.2518	1255.0209	1241.1319	1143.7861	1206.8658		

```
plastics_outlier <- as.xts(plastics)
plastics_outlier["000312"] <- 2023 #added 1000 to the original value
```


```
fit_outlier <- stl(as.ts(plastics_outlier), s.window="periodic", robust = T)
#Seasonally adjusted data from outlier data set
seasFit.outlier <- seasadj(fit_outlier)
seasFit.outlier
```

	Jan	Feb	Mar	Apr	May	Jun	Jul
## 1	993.6494	1014.5969	1013.6897	989.9799	983.6071	948.9579	965.7852
## 2	992.6494	1017.5969	1011.6897	1023.9799	1052.6071	1064.9579	1090.7852
## 3	1147.6494	1110.5969	1122.6897	1146.9799	1157.6071	1167.9579	1103.7852
## 4	1202.6494	1178.5969	1175.6897	1200.9799	1227.6071	1263.9579	1286.7852
## 5	1281.6494	1349.5969	1363.6897	1376.9799	1421.6071	1478.9579	1411.7852

	Aug	Sep	Oct	Nov	Dec
## 1	951.4535	933.5234	948.7386	993.7208	986.2975
## 2	1084.4535	1066.5234	1113.7386	1088.7208	1104.2975
## 3	1171.4535	1198.5234	1270.7386	1192.7208	2226.2975
## 4	1290.4535	1329.5234	1417.7386	1425.7208	1412.2975
## 5	1343.4535	1253.5234	1237.7386	1141.7208	1216.2975

The figure below shows the changes due to the outlier. We see the peak in the data and seasonally adjusted plots where the outlier exists. However, I cannot determine if there's any other effect on the graph due to the outlier.

```
autoplot(as.ts(plastics_outlier), series = "Data") +
  forecast::autolayer(seasFit.outlier, series="Seasonally Adjusted") +
  xlab("Year") + ylab("Sale Amount (in thousands)") +
  ggtitle("Product A monthly sales with seasonally adjusted") +
  scale_colour_manual(values=c("darkgoldenrod", "blue"),
    breaks=c("Data", "Seasonally Adjusted"))
```



Ch6.Q2.f)

Based on the values and the plot, it appears that there's no difference if the outlier is near the end rather than in the middle of the time series.

```
plastics_outlier <- as.xts(plastics)
plastics_outlier["000510"] <- 1920 #added 500 to the original value
fit_outlier_End <- stl(as.ts(plastics_outlier), s.window="periodic", robust = T)
seasFit.outlier.End <- seasadj(fit_outlier_End)
seasFit.outlier.End
```


	Jan	Feb	Mar	Apr	May	Jun	Jul
## 1	995.6289	1028.9663	1031.1597	1004.0011	984.6320	953.4195	970.5085
## 2	994.6289	1031.9663	1029.1597	1038.0011	1053.6320	1069.4195	1095.5085
## 3	1149.6289	1124.9663	1140.1597	1161.0011	1158.6320	1172.4195	1108.5085
## 4	1204.6289	1192.9663	1193.1597	1215.0011	1228.6320	1268.4195	1291.5085
## 5	1283.6289	1363.9663	1381.1597	1391.0011	1422.6320	1483.4195	1416.5085

	Aug	Sep	Oct	Nov	Dec
## 1	968.9721	938.6788	876.3292	997.7260	973.9780
## 2	1101.9721	1071.6788	1041.3292	1092.7260	1091.9780
## 3	1188.9721	1203.6788	1198.3292	1196.7260	1213.9780
## 4	1307.9721	1334.6788	1345.3292	1429.7260	1399.9780
## 5	1360.9721	1258.6788	1665.3292	1145.7260	1203.9780

Again, I cannot determine any difference in the graph below. Not sure, if it there's any effect on the graph due to the outlier apart from the peak at the outlier point.

```
autoplot(as.ts(plastics_outlier), series = "Data") +
  forecast::autolayer(seasFit.outlier.End, series="Seasonally Adjusted") +
  xlab("Year") + ylab("Sale Amount (in thousands)") +
  ggtitle("Product A monthly sales with seasonally adjusted") +
  scale_colour_manual(values=c("darkgoldenrod","blue"),
    breaks=c("Data","Seasonally Adjusted"))
```

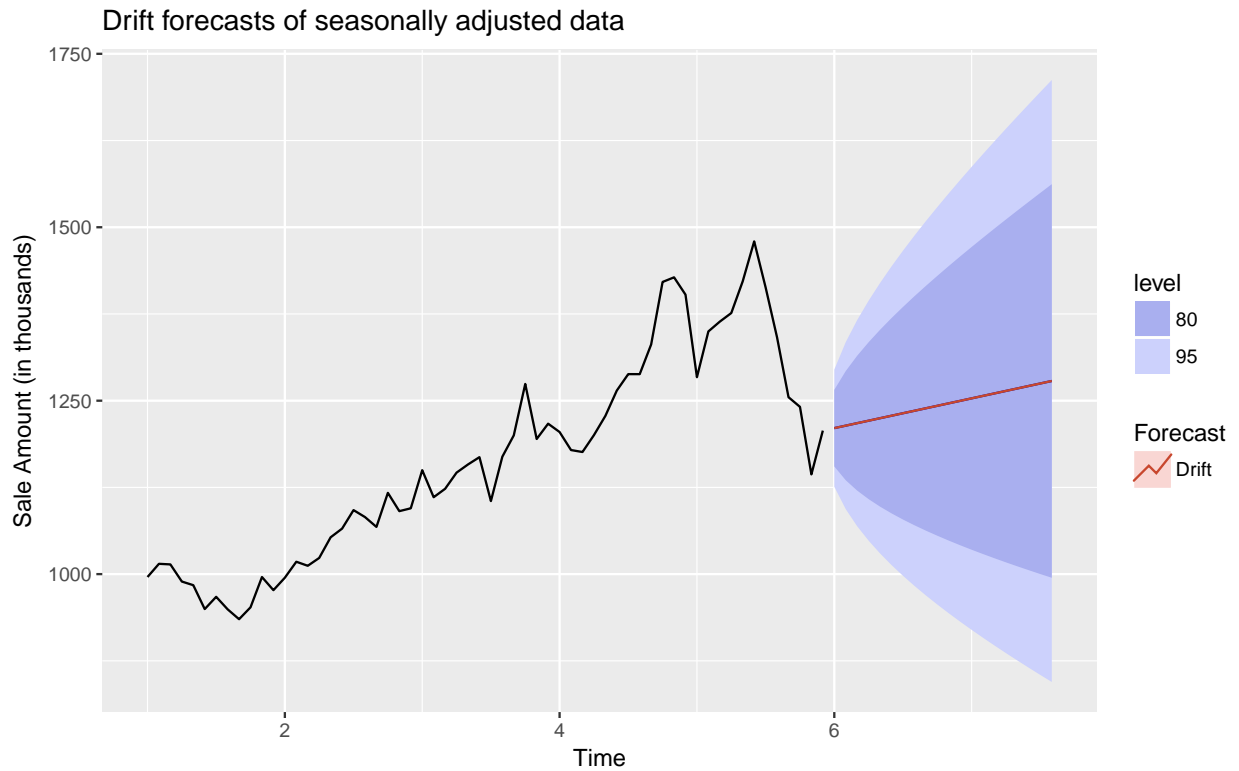
Product A monthly sales with seasonally adjusted



Ch6.Q2.g)

The figure below shows the forecast of the seasonally adjusted data using the random walk with drift.


```
autoplot(rwf(seasadj(fit), drift = TRUE, h=20)) +
  forecast::autolayer(rwf(seasadj(fit), drift = TRUE, h=20), PI=FALSE, series="Drift") +
  ylab("Sale Amount (in thousands)") +
  ggtitle("Drift forecasts of seasonally adjusted data") +
  guides(colour=guide_legend(title="Forecast"))
```



Ch6.Q2.h)

The figure below shows the forecast of the product A sales data based on the naive forecast of the seasonally adjusted data and a seasonal naive forecast of the seasonal component, after an STL decomposition of the data. The forecast is between 80% and 95% prediction intervals.

```
stlf(plastics, method = "naive") %>%
  autoplot() + ylab("Sale Amount (in thousands)")
```



Question 3

Ch6.Q3.a)

In Figure 6.13, the 'data' panel is the actual observed data. The three components ('seasonal', 'trend', and 'remainder') in the bottom panels can be added together to reconstruct the data shown in the 'data' panel. In the figure, we notice that the seasonal component changes very slowly over time. The trend-cycle is linear and increasing. the remainder component is shown in the bottom panel is what is left over when the seasonal and trend-cycle components have been subtracted from the data.

The large grey bar in the second panel shows that the variation in the seasonal component is small as compared to the variation in the data.

In Figure 6.13, the bottom graph shows the seasonal sub-series plot of the seasonal component. This kind of graph helps us to visualize the variation in the seasonal component over time. In this case, we notice there are only very small changes over time.

Ch6.Q3.b)

Yes, the recession of 1991/1992 is visible in the estimated components. This can be observed in the fourth panel i.e. remainder component. There is a negative peak in the remainder component for the year 1991 and 1992.